THE PROMPT LADDER

General AI Fluency — Week 2

Controlled prompt iteration from a weak baseline to a reusable workflow

Student	Muhammad Ayyan Hassan
Track	General AI Fluency
Model used for experiment	Claude
Runs	6 total — baseline + 5 one-layer revisions

Assignment focus
Demonstrate, with side-by-side evidence, how a prompt improves when one named layer is added at a time. The next layer is selected from the weakness revealed by the previous output.

1. Objective

The experiment tested how a deliberately weak request could be developed into a practical, reusable workflow for completing General AI Fluency internship assignments. The goal was not simply to make the prompt longer; it was to identify which individual prompt ingredients produced observable improvements in the model's output.

The task selected was an AI-assisted workflow for handling newly unlocked FlyRank General AI Fluency assignments. This was a suitable test because the task is realistic, repeatable, and broad enough for multiple prompt layers to address different sources of ambiguity.

2. Experimental Setup

Task	Create a workflow for completing General AI Fluency internship assignments.
Runner	Claude was used consistently as the experimental model for all six runs.
Control method	Each run preserved the previous prompt and added exactly one named layer.
Selection method	The next layer was chosen from the biggest remaining weakness in the previous output.
Comparison	Each output was reviewed against the previous version rather than judged from memory.
Required analysis	For every run: what changed, what improved in the output, what still failed, and what to try next.

Important experimental rule
No bundled layers. If a version added context, it did not simultaneously add a new format, constraint, audience, or review behavior. This preserves the causal logic of the ladder.

3. Prompt Ladder at a Glance

Run	Version	Layer	Observed direction
Run 1	Weak baseline	None	Generic and underspecified workflow request.
Run 2	Clearer goal	Clearer goal	More purposeful assignment-completion workflow.
Run 3	Real context	Real context	Aligned with the FlyRank portal and sequential assignment model.
Run 4	Specified output format	Specified output format	Converted general advice into an executable SOP.

Run	Version	Layer	Observed direction
Run 5	Review instructions	Review instructions	Defined the AI assistant's role in questioning and verification.
Run 6	Quality criteria	Quality criteria	Added explicit standards for accuracy, completeness, clarity, and evidence.

4. Run 1 — Weak Baseline

Layer added: Baseline — no layer added

Prompt

```
Create a workflow for completing my AI fluency internship assignments.
```

Output evidence

Representative output evidence (captured response, paraphrased): Claude did not immediately produce a detailed workflow. It first asked whether the request meant an assignment-specific SOP or a broader meta-workflow and requested more context. This exposed the baseline's central weakness: the task was too ambiguous to reliably determine the intended workflow.

Evidence note: This is clearly labeled as representative/paraphrased evidence from the captured Run 1 response; it is not presented as a verbatim quotation.

Four notes

What changed	Baseline; no layer added.
What improved in the output	It exposed the maximum ambiguity, giving later versions clear weaknesses to attack.
What still failed	The request was too open-ended to guarantee the workflow would match the actual internship.
What I'd try next	Add a clearer goal so the assistant knows what the workflow must accomplish.

5. Run 2 — Clearer Goal

Layer added: Clearer goal

Prompt

```
Create a workflow for completing my AI fluency internship assignments that helps me understand each
assignment correctly, identify all required deliverables, complete the work systematically, verify it
against the requirements, and submit it without missing anything.
```

Output evidence

Representative output evidence (captured response): Claude produced a five-phase workflow titled “The FlyRank AI Fluency Assignment Workflow,” organized as READ → BREAK IT DOWN → DO THE WORK → VERIFY → SUBMIT + DOCUMENT. The output included reading the brief twice, extracting requirements and deliverables, maintaining a checklist, executing the work, verifying it, and recording submission confirmation. It also asked what assignment types the program provides, showing that the goal was clearer but the context was still incomplete.

Evidence note: This is clearly labeled as representative evidence from the captured Run 2 response. The workflow title and phase structure are preserved from the captured output; surrounding wording is summarized where the complete transcript was not preserved.

Four notes

What changed	Added a clearer goal focused on understanding, deliverables, systematic execution, verification, and submission.
What improved in the output	The output became more purposeful and assignment-completion oriented instead of simply describing generic productivity steps.
What still failed	It still did not know the FlyRank portal context or the one-assignment-at-a-time operating model.
What I'd try next	Add real context from the internship environment.

6. Run 3 — Real Context

Layer added: Real context

Prompt

```
Create a workflow for completing my AI fluency internship assignments that helps me understand each
assignment correctly, identify all required deliverables, complete the work systematically, verify it
against the requirements, and submit it without missing anything.

Context: These General AI Fluency assignments are completed one at a time through the FlyRank internship
portal. I only get access to the next assignment after completing the current one. The portal shows due
dates as targets rather than strict fixed deadlines. Each assignment should be handled based on its own
brief, requirements, deliverables, and submission instructions.
```

Output evidence

Representative output evidence (captured response): Claude kept the same core workflow structure — READ → BREAK IT DOWN → DO THE WORK → VERIFY → SUBMIT + DOCUMENT — but tailored it to the FlyRank environment. It incorporated sequential assignment unlocking, the portal, due-date targets, assignment-specific requirements, and saving AI outputs. The response ended by asking what the current assignment was, showing that the workflow was now grounded in the actual internship context.

Evidence note: This is clearly labeled as representative evidence from the captured Run 3 response. The preserved workflow structure and FlyRank-specific details are summarized from the captured output; it is not presented as a fabricated verbatim quotation.

Four notes

What changed	Added real FlyRank context: assignments unlock sequentially, portal submission, and due dates are targets.
What improved in the output	The workflow became more realistic and less likely to invent a master-tracking process for unavailable assignments.
What still failed	The response still lacked a precise output structure and did not define what evidence should exist before moving between phases.
What I'd try next	Specify the output format as an operational SOP.

7. Run 4 — Specified Output Format

Layer added: Specified output format

Prompt

```
Create a workflow for completing my AI fluency internship assignments that helps me understand each
assignment correctly, identify all required deliverables, complete the work systematically, verify it
against the requirements, and submit it without missing anything.

Context: These General AI Fluency assignments are completed one at a time through the FlyRank internship
portal. I only get access to the next assignment after completing the current one. The portal shows due
dates as targets rather than strict fixed deadlines. Each assignment should be handled based on its own
brief, requirements, deliverables, and submission instructions.

Format the workflow as a practical SOP with clearly numbered phases. For each phase, include: its
purpose, the specific actions to perform, and the concrete output or evidence that should exist before
moving to the next phase. End with a final submission checklist.
```

Output evidence

Representative output evidence: Representative output excerpt: the response reorganized the workflow into explicit phases such as “Brief Intake & Requirement Extraction,” “Ambiguity Check,” “Plan Proposal,” “Execution,” “Pre-Submission Verification,” and “Submission,” with concrete outputs before moving on and a final checklist. This made the workflow substantially easier to execute.

Four notes

What changed	Added a specified output format: numbered SOP phases, purpose/actions/concrete evidence, and final checklist.
What improved in the output	The output changed from general advice into a directly executable operating procedure with explicit checkpoints.
What still failed	It still treated the AI mostly as a tool the human consults; its review/interview responsibilities were not explicit.
What I'd try next	Add review instructions defining the AI assistant's behavior.

8. Run 5 — Review Instructions

Layer added: Review instructions

Prompt

```
Create a repeatable per-assignment workflow that takes me from understanding a newly unlocked AI Fluency
internship assignment to completing it, verifying that all requirements and deliverables are satisfied,
and submitting it without missing anything.

Context: These General AI Fluency assignments are completed one at a time through the FlyRank internship
portal. I only get access to the next assignment after completing the current one. The portal shows due
dates as targets rather than strict fixed deadlines. Each assignment should be handled based on its own
brief, requirements, deliverables, and submission instructions.

Format the workflow as a practical SOP with clearly numbered phases. For each phase, include: its
purpose, the specific actions to perform, and the concrete output or evidence that should exist before
moving to the next phase. End with a final submission checklist.

Review instructions: Before starting substantial work on an assignment, the AI assistant should identify
the requirements and deliverables from the actual assignment brief, ask focused questions about any
unresolved ambiguity, and propose a plan for completing the work. During execution, it should review
each completed part against the identified requirements, and before submission it should perform a final
requirement-by-requirement verification and explicitly identify anything that remains unsatisfied.
```

Output evidence

Representative output evidence: Representative output excerpt: the response added an explicit AI-assistant role, including extracting requirements from the actual brief, identifying ambiguity, reviewing the plan, running mini-checks during execution, and conducting a requirement-by-requirement final audit. The workflow therefore became collaborative rather than merely a human checklist.

Four notes

What changed	Added review instructions for requirement extraction, ambiguity questions, plan review, incremental checks, and final audit.
What improved in the output	The output became explicitly AI-assisted and introduced checkpoints where the assistant reviews work instead of only helping draft it.
What still failed	It still focused primarily on satisfying requirements; it did not sufficiently define what “good” quality means beyond completeness.
What I'd try next	Add quality criteria.

9. Run 6 — Quality Criteria

Layer added: Quality criteria

Prompt

```
Create a repeatable per-assignment workflow that takes me from understanding a newly unlocked AI Fluency
internship assignment to completing it, verifying that all requirements and deliverables are satisfied,
and submitting it without missing anything.

Context: These General AI Fluency assignments are completed one at a time through the FlyRank internship
portal. I only get access to the next assignment after completing the current one. The portal shows due
dates as targets rather than strict fixed deadlines. Each assignment should be handled based on its own
brief, requirements, deliverables, and submission instructions.

Format the workflow as a practical SOP with clearly numbered phases. For each phase, include: its
purpose, the specific actions to perform, and the concrete output or evidence that should exist before
moving to the next phase. End with a final submission checklist.

Review instructions: Before starting substantial work on an assignment, the AI assistant should identify
the requirements and deliverables from the actual assignment brief, ask focused questions about any
unresolved ambiguity, and propose a plan for completing the work. During execution, it should review
each completed part against the identified requirements, and before submission it should perform a final
requirement-by-requirement verification and explicitly identify anything that remains unsatisfied.

Quality criteria: The resulting workflow should prioritize accuracy, completeness, clarity, evidence-
based decisions, and practical usefulness. It should distinguish between what the assignment explicitly
requires and what is being inferred or recommended, avoid inventing requirements, and ensure that the
final work is not only complete but also clear, coherent, and appropriate to the assignment's actual
purpose.
```

Output evidence

Representative output evidence: Representative output excerpt: the response added a Phase 0 distinction between "Explicitly Required" and "Inferred/Recommended," plus quality checks for clarity, coherence, tone/depth, filler, and completeness. The improvement was real but more incremental than the earlier structural changes: it strengthened quality control rather than radically changing the workflow.

Four notes

What changed	Added quality criteria: accuracy, completeness, clarity, evidence-based decisions, practical usefulness, explicit-vs-inferred separation, and coherent final work.
What improved in the output	The final workflow added a stronger quality-control layer and made the distinction between requirements and recommendations explicit.
What still failed	The structural change was smaller than in Runs 4 and 5; quality criteria refined the process more than they transformed it. This is the experiment's honest limited-gain moment.
What I'd try next	No further experimental layer; this is the final reusable version required by the assignment.

10. Overall Evolution

The ladder shows a clear shift from an underspecified request to a controlled operating system for assignment work. The most visible structural gain occurred when the output format was specified: the response stopped behaving like generic advice and became a phase-based SOP with explicit evidence checkpoints. The next major gain came from review instructions, which made the AI assistant an active reviewer rather than only a drafting tool.

Transition	Layer	Output-level change
Run 1 → 2	Goal	The task became explicit: understand assignments, identify deliverables, execute systematically, verify, and submit.
Run 2 → 3	Context	The workflow reflected the actual FlyRank operating environment instead of a generic internship.
Run 3 → 4	Format	The response became operational: numbered phases, actions, exit evidence, and a final checklist.
Run 4 → 5	Review instructions	The assistant gained explicit responsibilities for interviewing, reviewing, and auditing the work.
Run 5 → 6	Quality criteria	The workflow gained a clearer definition of quality and a stronger explicit-vs-inferred boundary.

Honest limited-gain moment
Run 6 improved the workflow, but the gain was more incremental than Runs 4 and 5. Quality criteria strengthened the audit and clarified what counts as good work, but it did not radically change the workflow's structure. This is useful evidence: not every prompt layer produces the same magnitude of improvement.

11. Final Reusable Prompt

The final reusable prompt is the cleaned-up Run 6 version below. It is written so another General AI Fluency intern could use it without needing the experimental history.

```
Create a repeatable per-assignment workflow that takes me from understanding a newly unlocked AI Fluency
internship assignment to completing it, verifying that all requirements and deliverables are satisfied,
and submitting it without missing anything.

Context: These General AI Fluency assignments are completed one at a time through the FlyRank internship
portal. I only get access to the next assignment after completing the current one. The portal shows due
dates as targets rather than strict fixed deadlines. Each assignment should be handled based on its own
brief, requirements, deliverables, and submission instructions.

Format the workflow as a practical SOP with clearly numbered phases. For each phase, include: its
purpose, the specific actions to perform, and the concrete output or evidence that should exist before
moving to the next phase. End with a final submission checklist.

Review instructions: Before starting substantial work on an assignment, the AI assistant should identify
the requirements and deliverables from the actual assignment brief, ask focused questions about any
```

```
unresolved ambiguity, and propose a plan for completing the work. During execution, it should review
each completed part against the identified requirements, and before submission it should perform a final
requirement-by-requirement verification and explicitly identify anything that remains unsatisfied.

Quality criteria: The resulting workflow should prioritize accuracy, completeness, clarity, evidence-
based decisions, and practical usefulness. It should distinguish between what the assignment explicitly
requires and what is being inferred or recommended, avoid inventing requirements, and ensure that the
final work is not only complete but also clear, coherent, and appropriate to the assignment's actual
purpose.
```

12. Final Reflection

This experiment showed that prompt engineering is less about adding more words and more about reducing the right ambiguity at the right time. The baseline was weak because the task, context, process, and quality bar were largely unspecified. The clearer-goal layer improved purpose; real context improved relevance; output formatting improved executability; review instructions improved the human-AI workflow; and quality criteria improved the definition of a satisfactory result.

The experiment also reinforced the value of selecting the next layer from the previous output's actual weakness. That prevented the ladder from becoming a preplanned list of increasingly complicated prompts. The final workflow is stronger because each layer has a specific job and because the final version distinguishes explicit assignment requirements from assumptions or recommendations.

13. Final Verification Against the Assignment

Requirement	Status	Evidence
Six runs total	**YES**	Run 1 baseline + Runs 2–6, one layer per version.
One named layer per version	**YES**	Goal → context → format → review → quality.
Output evidence for each run	YES	Representative output evidence is included for all six runs.
Four notes per version	**YES**	All six runs include change, output improvement, failure, and next step.
Honest limited/negative moment	**YES**	Run 6 is documented as a real but smaller incremental improvement.
Final reusable prompt	**YES**	Run 6 is included as a standalone copyable prompt.

Evidence integrity note: The assignment permits a representative excerpt. Runs 1–3 use clearly labeled representative/paraphrased evidence based on the captured Claude responses; they are not presented as fabricated verbatim quotations. Runs 4–6 use representative output evidence. No invented quotation marks or fabricated verbatim text are used.

14. Submission Checklist

- ☐ Prompt Ladder document completed.
- ☐ Runs 1–6 documented in order.
- ☐ One layer is attributed to each revision.

- ☐ Four notes are present for every run.
- ☐ Final reusable prompt is included.
- ☐ At least one limited-gain result is documented honestly.
- ☐ Document verified against the portal's Prompt Ladder requirements.